

Neural Network Model Report

Overview

I received a CSV containing more than 34,000 organizations that have received funding from Alphabet Soup over the years. Within this dataset are a number of columns that capture metadata about each organization. I was tasked with building a tool that can help it select the applicants for funding with the best chance of success in their ventures. With my knowledge of machine learning and neural networks, I used the features in the provided dataset to create a binary classifier that can predict whether applicants will be successful if funded by Alphabet Soup.

Data Processing

To process the data, I removed unnecessary columns ('EIN'). I then binned 'NAME', 'APPLICATION_TYPE', and 'CLASSIFICATION'. Once that was done, I use `pd.get_dummies` on the rest of the columns and set 'IS_SUCCESSFUL' to be the target for the model. The remaining columns were set to be features. The data was then split into training and testing and `StandardScaler` was used.

Compiling, Training, and Evaluating the Model

For my second attempt at evaluating the model, I used two hidden layers and one output layer. The first was a `relu` activation with 10 units. The second was another `relu` activation with 16 units. And the output layer was a 'sigmoid' activation with 1 unit.

I was able to achieve the target model performance, achieving a 79% accuracy. This is 7% higher than my first attempt, which achieved a 72% accuracy.

To increase the model's performance, I included the 'NAME' column, which was previously dropped from the database, and binned it along with 'APPLICATION_TYPE', and 'CLASSIFICATION'. I played around with the number of hidden layers and units to use. Ultimately, I settled on using three more units on my first hidden layer, and two more units on my second hidden layer than I did previously on my first attempt.

First Attempt:

Model: "sequential_1"

Layer (type)	Output Shape	Param #
dense_3 (Dense)	(None, 7)	308
dense_4 (Dense)	(None, 14)	112
dense_5 (Dense)	(None, 1)	15

=====
Total params: 435
Trainable params: 435
Non-trainable params: 0

268/268 - 0s - loss: 0.5557 - accuracy: 0.7256 - 374ms/epoch - 1ms/step
Loss: 0.5557180643081665, Accuracy: 0.7255976796150208

Second Attempt:

Model: "sequential_9"

Layer (type)	Output Shape	Param #
dense_32 (Dense)	(None, 10)	4470
dense_33 (Dense)	(None, 16)	176
dense_34 (Dense)	(None, 1)	17

=====
Total params: 4,663
Trainable params: 4,663
Non-trainable params: 0

268/268 - 0s - loss: 0.4474 - accuracy: 0.7930 - 474ms/epoch - 2ms/step
Loss: 0.44742581248283386, Accuracy: 0.7930029034614563

Summary

In conclusion, the second model performed better all-around at a 79% accuracy and 0.45 loss. Whereas the first model had a 72% accuracy with a 0.55 loss. My recommendation would be to use the second model. In the future I would use an automated neural network to determine the best model hyperparameters to use.